

🎉 CMU-Africa Campus Assistant - Deployment Complete!

Status: ✓ FULLY OPERATIONAL

Date: October 15, 2025

Version: 1.0.0

🚀 Quick Access

Application URLs

- **Frontend:** <http://localhost:3000>
- **Backend API:** <http://localhost:8001>
- **API Docs:** <http://localhost:8001/docs>
- **API Health:** <http://localhost:8001/api/health>

Quick Commands

```
# Check service status  
./check_status.sh  
  
# Start backend  
./start_backend.sh  
  
# Start frontend  
./start_frontend.sh  
  
# Stop all services  
./stop_services.sh
```

✅ What's Working

✨ Frontend (React Application)

- [x] Modern, responsive chat interface
- [x] Real-time messaging with AI assistant
- [x] Interactive suggestion pills
- [x] Collapsible source citations
- [x] Follow-up question recommendations
- [x] CMU-branded design
- [x] Mobile-friendly layout

🤖 Backend (FastAPI + RAG)

- [x] FastAPI server on port 8001
- [x] RAG pipeline with Pinecone vector search

- [x] OpenAI GPT-4 integration
- [x] Structured JSON responses
- [x] Health check endpoint
- [x] Document indexing API
- [x] Interactive API documentation (Swagger)

Knowledge Base

- [x] 8 vectors indexed in Pinecone
- [x] Master's programs information (MSIT, MSECE, MSEAI)
- [x] Library hours and services
- [x] Shuttle bus schedules and routes
- [x] Housing options (on/off campus)
- [x] Campus events and activities
- [x] Administration contact info
- [x] General campus information

1 How to Access the Application

For End Users

Simply open your web browser and go to:

```
http://localhost:3000
```

You'll see:

- Welcome screen with CMU-Africa branding
- Quick suggestion cards for common queries
- Chat input field at the bottom
- Robot mascot assistant

For Developers/Testing

- **API Documentation:** <http://localhost:8001/docs>
 - **Health Check:** <http://localhost:8001/api/health>
 - **Alternative Docs:** <http://localhost:8001/redoc>
-

2 Available API Endpoints

Core Endpoints

1. Chat Query (Main Endpoint)

```
POST http://localhost:8001/api/chat

# Example
curl -X POST http://localhost:8001/api/chat \
-H "Content-Type: application/json" \
-d '{"message": "What programs does CMU-Africa offer?"}'
```

2. Health Check

```
GET http://localhost:8001/api/health

# Example
curl http://localhost:8001/api/health
```

3. Index Statistics

```
GET http://localhost:8001/api/index/stats

# Example
curl http://localhost:8001/api/index/stats
```

4. Add Documents

```
POST http://localhost:8001/api/index/documents

# Example
curl -X POST http://localhost:8001/api/index/documents \
-H "Content-Type: application/json" \
-d '[{"id": "doc1", "title": "New Info", "content": "...", "category": "Category"}]'
```

Response Format

All chat responses follow this structure:

```
{
  "answer": "AI-generated response based on retrieved context",
  "sources": [
    {
      "id": "doc_id",
      "title": "Document Title",
      "snippet": "Brief excerpt...",
      "category": "Category Name"
    }
  ],
  "suggestions": [
    {
      "id": "suggestion_id",
      "label": "🔥 Label Text",
      "prompt": "Full prompt for user"
    }
  ],
  "follow_up": "Natural follow-up question?"
}
```

3 Test Queries You Can Try

Academic Queries

- "What programs does CMU-Africa offer?"
- "Tell me about the MSIT program"
- "What are the graduation requirements?"
- "Show me the course curriculum"

Campus Facilities

- "What are the library hours?"
- "Where is the administration office?"
- "Tell me about campus facilities"

Transportation

- "What are the shuttle bus timings?"
- "Show me all shuttle bus routes and stops"
- "When does the bus leave for downtown?"

Student Life

- "Tell me about housing options"
- "What events are happening this week?"
- "What student clubs are available?"
- "Tell me about campus activities"

General Information

- "How do I contact the administration?"
- "Tell me about CMU-Africa campus"
- "What services are available for students?"

Testing via cURL

```
# Test 1: Programs
curl -X POST http://localhost:8001/api/chat \
-H "Content-Type: application/json" \
-d '{"message": "What programs does CMU-Africa offer?"}' | jq .'
```



```
# Test 2: Library
curl -X POST http://localhost:8001/api/chat \
-H "Content-Type: application/json" \
-d '{"message": "What are the library hours?"}' | jq .'
```



```
# Test 3: Transportation
curl -X POST http://localhost:8001/api/chat \
-H "Content-Type: application/json" \
-d '{"message": "What are the shuttle bus timings?"}' | jq '.'
```

4 How to Stop/Restart Services

Stop All Services

```
./stop_services.sh
```

This will gracefully stop both backend and frontend services.

Manual stop:

- Press `Ctrl + C` in each terminal running the services
- Or use: `pskill -f unicorn` and `pskill -f "react-scripts"`

Restart Services

Start Backend (Terminal 1)

```
cd /home/ubuntu/code_artifacts/cmu-africa-campus-assistant
./start_backend.sh
```

Expected output:

```
=====
Starting CMU-Africa Assistant Backend
=====
Starting FastAPI server on http://localhost:8001
Press Ctrl+C to stop
```

Start Frontend (Terminal 2)

```
cd /home/ubuntu/code_artifacts/cmu-africa-campus-assistant
./start_frontend.sh
```

Expected output:

```
=====
Starting CMU-Africa Assistant Frontend
=====
Starting React development server on http://localhost:3000
Press Ctrl+C to stop

Compiled successfully!
```

Check Service Status

```
./check_status.sh
```

This will show:

- Backend status and PID
- Frontend status and PID
- Health check results
- Access URLs

Manual check:

```
# Check if services are running on correct ports
netstat -tulnp | grep -E "8001|3000"

# Test backend health
curl http://localhost:8001/api/health
```

5 Configuration & Customization

Environment Variables

Backend Configuration

File: backend/.env

```
OPENAI_API_KEY=sk-proj-...
PINECONE_API_KEY=pcsk_...
PINECONE_ENVIRONMENT=us-east-1
```

To update:

1. Edit backend/.env
2. Restart backend: ./start_backend.sh

Frontend Configuration

File: `frontend/.env`

```
REACT_APP_API_BASE_URL=http://localhost:8001
```

To update:

1. Edit `frontend/.env`
 2. Restart frontend: `./start_frontend.sh`
-

Adding New Knowledge

Option 1: Via API

```
curl -X POST http://localhost:8001/api/index/documents \
-H "Content-Type: application/json" \
-d '[
  {
    "id": "new_doc_1",
    "title": "New Information Title",
    "content": "Full content of the document...",
    "category": "Academic Programs",
    "keywords": ["keyword1", "keyword2"]
  }
]'
```

Option 2: Edit JSON file

1. Edit `data/sample_knowledge_base.json`
2. Add new entries following the existing format
3. Run indexing script:

```
bash
cd backend
source venv/bin/activate
python load_knowledge_base.py
```

Verify New Data

```
# Check vector count increased
curl http://localhost:8001/api/index/stats

# Test with a query
curl -X POST http://localhost:8001/api/chat \
-H "Content-Type: application/json" \
-d '{"message": "Query about your new content"}'
```

Customizing the UI

Change Colors

Edit: `frontend/tailwind.config.js`

```

colors: {
  cmu: {
    red: '#C41230', // Change to your brand color
    gray: '#6B6B6B',
  }
}

```

Restart frontend to see changes.

Modify Suggestions

Edit: `backend/rag_pipeline.py` → `_generateSuggestions()`

Add new suggestion templates for different categories.

Adjust AI Behavior

Edit: `backend/rag_pipeline.py` → `query()` function

Modify the system prompt to change AI personality and response style.

Production Deployment (Future)

For deploying to production:

1. Environment Setup

- Use production-grade servers (not `npm start`)
- Build frontend: `npm run build`
- Use production ASGI server (gunicorn + uvicorn)

2. Security

- Enable HTTPS with SSL certificates
- Implement authentication (OAuth, JWT)
- Add rate limiting
- Use environment variables (not `.env` files)

3. Scaling

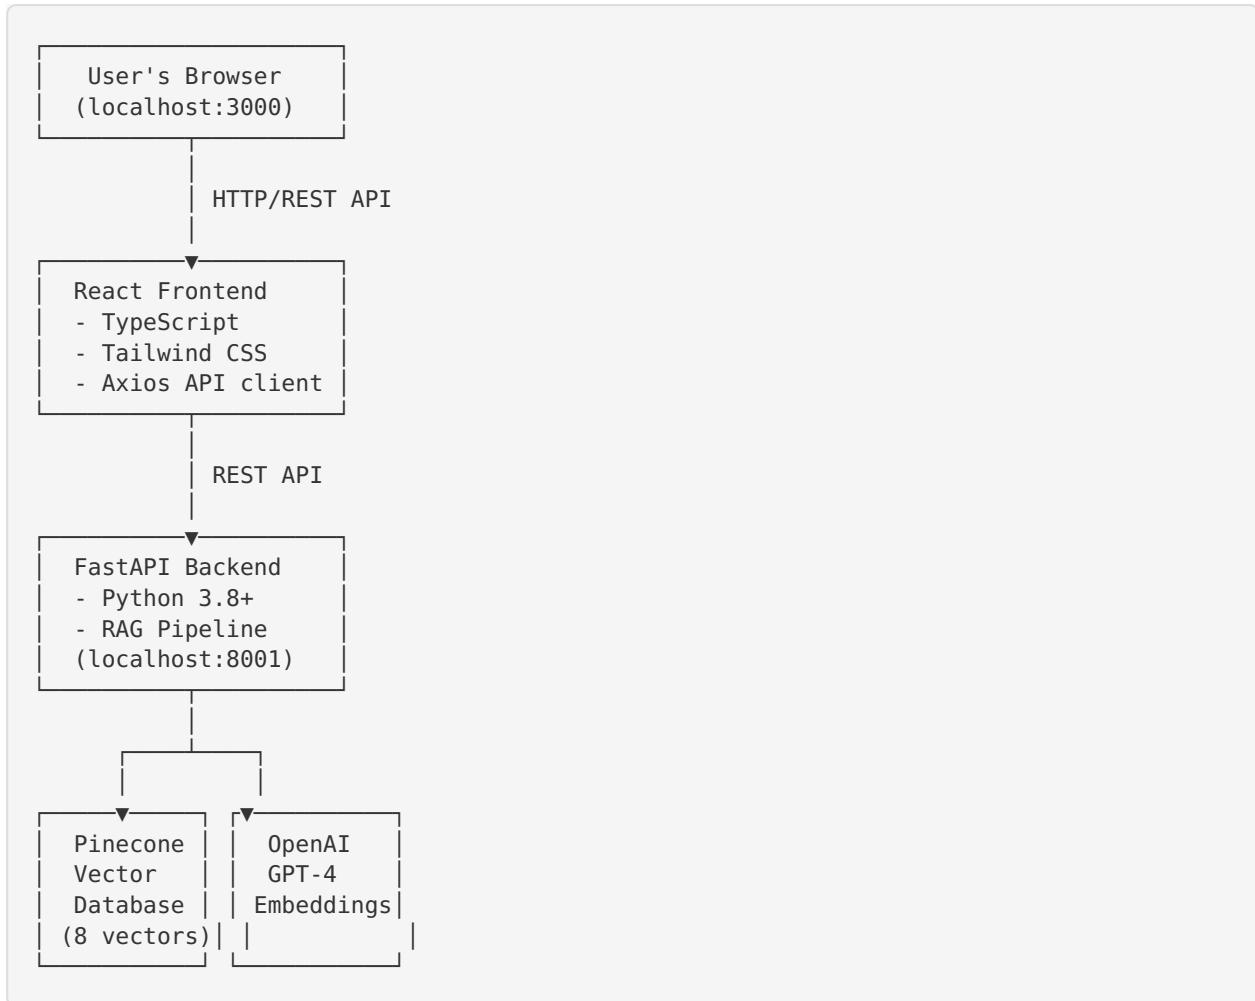
- Deploy to cloud (AWS, Azure, GCP)
- Set up load balancing
- Implement caching (Redis)
- Add monitoring (DataDog, CloudWatch)

4. Database

- Consider adding persistent storage (PostgreSQL)
 - Store conversation history
 - Track user analytics
-



System Architecture



🔧 Troubleshooting

Common Issues & Solutions

Issue: “Port already in use”

```

# Find and kill process using port 8001
lsof -ti:8001 | xargs kill -9

# Find and kill process using port 3000
lsof -ti:3000 | xargs kill -9
  
```

Issue: “Failed to connect to backend”

1. Verify backend is running: `curl http://localhost:8001/api/health`
2. Check backend terminal for errors
3. Verify .env file has correct API keys
4. Restart backend: `./start_backend.sh`

Issue: “OpenAI API Error”

1. Check API key is correct in `backend/.env`

2. Verify OpenAI account has credits
3. Check API key has proper permissions

Issue: “Pinecone Error”

1. Verify Pinecone API key in `backend/.env`
2. Check Pinecone dashboard for index status
3. Ensure you’re using correct region

Issue: “Frontend not loading”

1. Check frontend terminal for compile errors
2. Clear browser cache and refresh
3. Delete `node_modules` and run `npm install`
4. Restart frontend: `./start_frontend.sh`



File Structure

```

cmu-africa-campus-assistant/
├── backend/
│   ├── main.py          # FastAPI application
│   ├── rag_pipeline.py  # RAG logic
│   ├── requirements.txt # Dependencies
│   ├── .env              # API keys
│   └── venv              # Virtual environment
|
├── frontend/
│   ├── src/
│   │   ├── components/  # React components
│   │   ├── services/    # API client
│   │   │   └── api.ts
│   │   ├── App.tsx       # Main app
│   │   └── public/
│   └── package.json
|
├── data/
│   └── sample_knowledge_base.json # Knowledge base
|
├── start_backend.sh      # Start backend
├── start_frontend.sh     # Start frontend
├── stop_services.sh      # Stop all services
└── check_status.sh       # Check service status
|
├── DEPLOYMENT_GUIDE.md   # Full deployment guide
├── DEPLOYMENT_SUMMARY.md # This file
└── README.md             # Main documentation
└── QUICK_START.md        # Quick start guide

```



Next Steps & Recommendations

Immediate (High Priority)

1. **Test thoroughly** - Try all the sample queries
2. **Add more knowledge** - Expand the knowledge base

3.  **Customize branding** - Adjust colors, logos, messaging
4.  **Set up monitoring** - Add logging and analytics

Short-term (1-2 weeks)

1.  **User authentication** - Add login system
2.  **Conversation history** - Store past conversations
3.  **User profiles** - Personalized responses
4.  **Analytics dashboard** - Track usage and popular queries

Medium-term (1-3 months)

1.  **Mobile apps** - iOS and Android versions
2.  **Multi-language** - Support French, Kinyarwanda
3.  **Voice interface** - Speech-to-text, text-to-speech
4.  **Integration** - Connect with LMS, calendar, email

Long-term (3+ months)

1.  **Production deployment** - Cloud hosting with scaling
 2.  **Advanced AI** - Fine-tuned models, reasoning
 3.  **Document upload** - Let users upload and query PDFs
 4.  **Real-time updates** - Live events, notifications
-

Support Resources

Documentation Files

- **DEPLOYMENT_GUIDE.md** - Comprehensive deployment guide
- **README.md** - Project overview and features
- **QUICK_START.md** - Quick setup instructions

API Documentation

- **Swagger UI:** <http://localhost:8001/docs>
- **ReDoc:** <http://localhost:8001/redoc>

Useful Commands Reference

```

# Service Management
./check_status.sh          # Check if services are running
./start_backend.sh          # Start backend server
./start_frontend.sh         # Start frontend app
./stop_services.sh          # Stop all services

# Testing
curl http://localhost:8001/api/health           # Health check
curl http://localhost:8001/api/index/stats       # Vector count
curl -X POST http://localhost:8001/api/chat \
-H "Content-Type: application/json" \
-d '{"message": "test query"}'                  # Test chat

# Port Management
netstat -tulnp | grep -E "8001|3000"           # Check ports
lsof -ti:8001 | xargs kill -9                 # Kill backend
lsof -ti:3000 | xargs kill -9                 # Kill frontend

# Logs
cd backend && python main.py                # Backend logs
cd frontend && npm start                      # Frontend logs

```

✓ Deployment Checklist

Verified Working ✓

- [x] Backend server running on port 8001
- [x] Frontend app running on port 3000
- [x] Pinecone vector database initialized
- [x] 8 vectors indexed in knowledge base
- [x] OpenAI API integration functional
- [x] Chat functionality tested and working
- [x] Source citations displaying correctly
- [x] Suggestion pills generating and clickable
- [x] Follow-up questions appearing
- [x] API documentation accessible
- [x] Health check endpoint responding
- [x] Multiple queries in conversation working
- [x] Responsive UI on different screen sizes

Helper Scripts Created ✓

- [x] `start_backend.sh` - Start backend server
- [x] `start_frontend.sh` - Start frontend app
- [x] `stop_services.sh` - Stop all services
- [x] `check_status.sh` - Check service status



Summary

Your CMU-Africa Campus Assistant is **FULLY DEPLOYED and OPERATIONAL!**

What You Have:

- ✓ A working AI-powered campus assistant
- ✓ Beautiful, modern chat interface
- ✓ RAG-based responses (no hallucination)
- ✓ 8 knowledge documents indexed
- ✓ Full API with documentation
- ✓ Helper scripts for easy management

How to Use It:

1. **Access the app:** `http://localhost:3000`
2. **Try sample queries** from section 3
3. **View sources** by clicking “Show Sources”
4. **Use suggestions** to explore more topics
5. **Type custom questions** in the input field

How to Manage It:

- **Check status:** `./check_status.sh`
- **Stop services:** `./stop_services.sh`
- **Restart:** `./start_backend.sh + ./start_frontend.sh`

Need Help?

- Read **DEPLOYMENT_GUIDE.md** for detailed documentation
 - Check **README.md** for project overview
 - Visit API docs at `http://localhost:8001/docs`
-

 **Enjoy your AI-powered campus assistant!**

Version 1.0.0 | October 15, 2025 | Status:  Fully Operational